

---

# THE ESTIMATOR DECIDES: WHAT CURRICULA AND FAILURE RECYCLING CAN AND CANNOT DO IN RL WITH VERIFIABLE REWARDS

Anonymous authors

Paper under double-blind review

## ABSTRACT

Post-training with verifiable rewards increasingly relies on two data-level interventions: *difficulty curricula*, which reallocate rollout budget toward learnable prompts, and *failure recycling* (hindsight relabeling), which converts failed rollouts into verified successes of the tasks they actually achieved. Both are treated as objective-agnostic add-ons. We show they are not: the advantage estimator underneath shapes what each can contribute, in closed form. For success-conditioned (MaxRL-style) advantages, the expected absolute coefficient mass a prompt emits from  $N$  rollouts is exactly  $2(\text{pass}@N - \text{pass}@1)$  — a compute-indexed learnability functional, vanishing exactly at pass rates 0 and 1 with a single interior peak set by the rollout budget alone — and the practical drop-all-fail implementation is unbiased for the truncation-order  $T=N-1$  objective, not  $T=N$ . Coefficient mass is a surrogate (it is not gradient norm), but it is knowable before spending a rollout, and it predicts an estimator-conditioned divergence we then measure: across nine heterogeneous maze runs, every success-conditioned run grew  $\text{pass}@k$  coverage and every GRPO run lost it — consistent, exploratory evidence that motivates the balanced factorial we outline. A two-seed GSM8K pilot points the same direction (teacher-deficit sign under GRPO in 2/2 seeds; the registered regression shape in 1 of 2). The same algebra flags a failure mode of recycling that we observe and replicate at three seeds on Countdown: relabels to already-reachable destinations improve mean accuracy while *losing*  $\text{pass}@k$  coverage — **recycling-induced sharpening** — and it motivates the mitigation, gating relabel admission by the destination’s estimated pass rate. At the frontier tier the moderate gate returned coverage to the no-recycling baseline while keeping most of the mean gain; the tested settings give one useful operating point, not yet a monotonicity guarantee. The boundary of the safety claim is measured too: a pre-registered fifth suite whose learnable region collapses to one template stratum defeats *every* estimator. One diagnosis covers all of it: audit the estimator before shipping a curriculum or a recycler, and measure coverage alongside mean accuracy — the currency in which every failure mode here is visible.

## 1 INTRODUCTION

Post-training language models and control policies with verifiable rewards has a cost structure unlike supervised learning: the data is free, but *rollouts* are not. Every optimization step is preceded by sampling  $N$  completions per task, generation can dominate wall-clock (in our committed telemetry, eight generation-dominated sessions spanning 282 steps average 83% of step time; across all 1,188 mixed GSM8K/Countdown steps, including CPU-bottlenecked maze-scale sessions, the step-weighted mean is 31%), and on hard task pools most of that compute buys nothing — under uniform sampling we measure that **65–75% of rollout groups produce zero learning signal**: either every rollout fails (the group is dropped by success-conditioned estimators) or every rollout succeeds (no contrast, no gradient).

The field’s response is a pair of data-level interventions. The first is the *difficulty curriculum*: reallocate the rollout budget toward tasks where learning is possible right now — UCB bandits over difficulty buckets (Wang et al., 2025), target-difficulty controllers (Shi et al., 2025), category bandits (Chen et al., 2025), learnability-based rejection sampling (Foster et al., 2025), dead-prompt filtering (Yu et al., 2025; Zheng et al., 2025). The second, newer, is *failure recycling*: convert a failed rollout into a verified success of the task it actually achieved — hindsight relabeling, recently brought into RLVR-style loops for robot policies (Xu et al., 2026). Both interventions are, in current practice, designed and evaluated as if they were *objective-agnostic*: bolt-on modules whose benefit is assumed to transfer across the advantage estimator they sit on, and whose cost is assumed visible in the mean-accuracy metrics they are tuned against.

This paper shows that neither assumption survives contact with the estimator’s own algebra. The organizing fact is the shape of one function. **The estimator’s expected coefficient mass  $u_N(p)$  vanishes exactly at pass rates 0 and 1 and peaks once in between, tracing three regimes of difficulty: a dead zone, a learnable region, and a mastered tail. Curricula reallocate compute within the learnable region; verified relabeling is a way to create a positive update for an all-fail requested task; and relabels that land in the mastered tail buy sharpening, not signal. The estimator shapes all three.** Three questions unpack this:

*Q1: What can a curriculum contribute, at most?* For success-conditioned advantages, the expected coefficient mass a task emits from a group of  $N$  rollouts has a closed form:  $2(\text{pass}@N - \text{pass}@1)$ , a compute-indexed learnability functional whose two exact zeros anchor the difficulty regimes (§3). Mass is a surrogate for learning signal (Remark 1), but it is the part a sampler can act on before spending a rollout, and it makes the allocation *derivable* rather than heuristic: the learnable region’s location, peak, and  $N$ -scaling come from the budget you already chose, and published learnability curricula are the  $N=2$  slice. It also bounds the channel. Allocation can only redistribute signal that exists: a perfect-information allocator ties our cheap posterior in end performance (§6.1), and at LLM scale a prompt-level teacher starves before it pays (§6).

*Q2: When is either intervention safe?* The same algebra says the answer depends on the estimator. MaxRL concentrates  $(N-1) \times$  more mass than RLOO on hard tasks as  $p \rightarrow 0$  and releases mastered tasks; GRPO’s variance-normalized weights keep spending on both tails. We predict and measure the resulting divergence at two scales: across the audited maze runs, every MaxRL-labeled run grows pass@8 coverage and every GRPO-labeled run loses it (§6.3; exploratory — the run pool is heterogeneous); on GSM8K, the pre-registered prediction that GRPO+teacher would be the only regressing cell landed on the registered run (§6.7). Recycling has its own failure mode, and the algebra flags it: relabels that land in the mastered tail sharpen. On Countdown arithmetic this is **recycling-induced sharpening** — exact-verifier relabeling improves mean accuracy while *losing* pass@ $k$  coverage — the data-induced sibling of GRPO’s weight-induced collapse, formally hindsight’s marginal-distribution shift (Ghosh et al., 2021) composed with curation-driven self-consumption (Ferbach et al., 2024).

*Q3: Does the algebra also supply the fix?* It motivates one. A relabel to a destination the model reliably reaches sits in the mastered tail, where mass vanishes, so we gate relabel admission by the destination’s estimated pass rate: the high- $p$  zero of the same utility that schedules sampling, read as a saturation threshold on what the recycler may admit. Across three seeds the moderate gate returns frontier-tier coverage to the no-recycling baseline while keeping most of recycling’s mean gain; the available settings (moderate at three seeds, strong at one seed under corrected code) sketch a mean-versus-coverage trade but do not yet establish a monotone dial (§6.9). Where recycling is healthy the gate is transparent: hard-destination relabels pass untouched, so the validated testbeds reproduce their results with the gate enabled. One method, no per-domain switches.

Our contributions:

1. **A closed-form curriculum (§3)**: the expected coefficient mass of a prompt under success-conditioned RL is exactly  $2(\text{pass}@N - \text{pass}@1)$ , a compute-indexed learnability functional whose peak location and  $N$ -scaling are derived rather than tuned, and whose exact zeros at  $p \in \{0, 1\}$  anchor the dead zone, the learnable region, and the mastered tail; published learnability curricula are its  $N=2$  slice.
2. **A derived teacher and a characterized recycler (§4–5)**: FrontierMax schedules prompts by this functional with the per-group estimator untouched, and hindsight relabeling under a verified-




Figure 1: **The estimator is the curriculum.** *Left:* the derived utility  $u_N(p) = (1 - (1 - p)^N) - p$  — half the estimator’s expected coefficient mass  $A_N = 2u_N$  (Prop. 1) — for growing rollout budgets  $N$ . The peak  $p^* \approx \ln N/N$  walks toward harder prompts as compute grows; the utility vanishes exactly at  $p \in \{0, 1\}$ . *Right:* all three curves are exact expected half-mass at  $N=16$  for the estimators as deployed (GRPO with sample-SD normalization:  $\sqrt{\frac{N-1}{N}} \frac{1}{N} \mathbb{E} \sqrt{K(N-K)}$ ,  $K \sim \text{Bin}(N, p)$ ; MC-verified against the code). MaxRL concentrates  $\rightarrow (N-1) \times$  more mass than RLOO on frontier prompts (Prop. 2) and releases mastered prompts fastest; GRPO’s variance normalization keeps  $\sim (N-1)/\sqrt{N} \times$  MaxRL’s relative mass on mastered prompts — the maintenance-of-easy-tasks its curriculum arms lose in §6. RLOO’s profile is exactly the published “learnability” objective — the  $N=2$  slice.

success contract is characterized as a verified, adaptively selected auxiliary update — on-policy for the achieved task only under an explicit distribution-matching condition (Remark 3).

3. **An estimator-conditioned coverage divergence (§6):** under identical curricula, success-conditioned runs grow  $\text{pass}@k$  coverage while GRPO runs lose it — zero exceptions across the audited maze cohort, with the same sign at LLM scale in a two-seed pilot; we grade this evidence as exploratory-but-consistent and specify the balanced factorial that would confirm it.
4. **A new failure mode and its derived mitigation (§6.8–6.9):** exact-verifier recycling improves mean accuracy while losing coverage — recycling-induced sharpening, replicated at three seeds — and gating relabels by destination pass rate yields an operating point that restores frontier-tier coverage while keeping most of the mean gain.

Along the way we state a refinement to the MaxRL paper — the practical drop- $K=0$  estimator is unbiased for the  $T=N-1$  objective, not  $T=N$  (Lemma 1) — and every prediction below was committed before the deciding runs, with negatives and retractions reported in place (§8) and audited in the repository.

## 2 BACKGROUND: MAXRL IN THREE LINES

For binary rewards, maximum likelihood maximizes  $\mathbb{E}[\log p_\theta(\text{success} \mid x)]$ . MaxRL expands  $\log p$  as a Maclaurin series in  $(1 - p)$  and truncates at order  $T$ , giving a compute-indexed family  $J^{(T)}$  interpolating REINFORCE ( $T=1$ ) to exact ML ( $T \rightarrow \infty$ ); the practical estimator sets per-rollout advantages  $w_i = r_i/K - 1/N$  with all-fail groups dropped (Lemma 1 gives its exact truncation order). Its weight function grows as  $p \rightarrow 0$ : the objective is an *implicit, gradient-level* curriculum. Our work is the data-level complement: the same algebra, read as a sampling rule plus a recycling rule.

## 3 THE ESTIMATOR IS THE CURRICULUM

**Lemma 1** (Truncation order of the practical estimator). *Let  $J_T(p) = -\sum_{t=1}^T (1 - p)^t / t$  denote the  $T$ -truncated ML objective, so  $\nabla J_T = (\sum_{j=0}^{T-1} (1 - p)^j) \nabla p$ ;  $T \rightarrow \infty$  recovers exact ML and  $T=1$  REINFORCE. Under  $N$  i.i.d. binary rewards with pass rate  $p$ , the drop- $K=0$  variant of the practical*

MaxRL estimator has expectation  $(\frac{1-(1-p)^N}{p} - (1-p)^{N-1})\nabla p = (\sum_{j=0}^{N-2} (1-p)^j)\nabla p$ : it is exactly unbiased for  $J_{N-1}$ , not  $J_N$ , for  $N \geq 2$ . Zeroing the control variate only on all-fail groups makes it outcome-dependent, contributing the  $-(1-p)^{N-1}\nabla p$  correction, and  $\sum_{k=1}^N (1-p)^{k-1} - (1-p)^{N-1} = \sum_{k=1}^{N-1} (1-p)^{k-1}$ .

**Interpretation.** A refinement to [Tajwar et al. \(2026\)](#), immaterial for their conclusions but load-bearing for our mass accounting at small  $N$  — all identities below are for the estimator as actually shipped.

**Proposition 1** (Expected coefficient mass). *For a prompt with pass rate  $p$  and  $N$  i.i.d. rollouts under the practical MaxRL weights, define the utility  $u_N(p) := \text{pass@}N(p) - \text{pass@}1(p) = (1 - (1-p)^N) - p$ . Then the expected absolute coefficient mass is*

$$A_N(p) := \mathbb{E}\left[\sum_i |w_i|\right] = 2(\text{pass@}N(p) - \text{pass@}1(p)) = 2u_N(p). \quad (1)$$

*Proof (three lines): on  $\{K \geq 1\}$ ,  $\sum_i |w_i| = K(\frac{1}{K} - \frac{1}{N}) + \frac{N-K}{N} = 2(1 - \frac{K}{N})$ , and  $\mathbb{E}[2(1 - \frac{K}{N})\mathbf{1}\{K \geq 1\}] = 2(\Pr[K \geq 1] - p) = 2(\text{pass@}N - \text{pass@}1)$ . The sampling utility used everywhere below (figures, Algorithm 1, code) is  $u_N$ ; the factor two matters only when quoting exact mass, and we carry it explicitly as  $A_N = 2u_N$ . A useful exact companion:  $u_N(p) = p(1-p)w_{N-1}(p)$  with  $w_T(p) = (1 - (1-p)^T)/p$  the truncated weight function of Lemma 1 — mass and the practical estimator’s weight view are linked, not identical objects.*

**Interpretation.** The estimator’s expected coefficient mass on a prompt is twice the probability that the prompt is solvable within  $N$  attempts but not within one. Coefficient mass is a surrogate for learning signal (Remark 1), but its two zeros are exact, and they anchor three regimes of difficulty: a dead zone at  $p = 0$  (nothing to condition on; under the current policy and support, sampling alone cannot produce a success there — and at small  $p > 0$  a task is operationally dead at budget  $(B, N)$  when its live-group probability  $1 - (1 - \ell_N(p))^B$ , with  $\ell_N(p) = 1 - p^N - (1-p)^N$  and  $B$  its share of groups, is too small to sustain training: on the frontier-heavy pool of \$6.2 ( $p \leq 10^{-5}$ ,  $\sim 133$  groups per task) that probability is  $\leq 2\%$  per task, and the one live group observed across all uniform seeds did not ignite — dead in the budgeted sense, not by fiat), a learnable region with a unique interior peak at  $p^* = 1 - N^{-1/(N-1)} \approx \ln N/N$  (unique because  $u_N'' = -N(N-1)(1-p)^{N-2} < 0$ ), and a mastered tail at  $p = 1$  (no contrast left). The regimes are continuous — the only exact zeros are  $p \in \{0, 1\}$  — so where we quote a “band” we mean an explicit threshold set  $\{p : u_N(p) \geq \eta u_N(p^*)\}$  ( $\eta=0.5$  unless stated), whose location and  $N$ -scaling follow from the one parameter you already chose. Everything in this paper is one of three moves against this structure: allocate within the learnable region (§4), create signal at its bottom (§5), and refuse updates that land in the mastered tail (the gate).

**Proposition 2** (Frontier amplification). *As  $p \rightarrow 0$ , the ratio of MaxRL’s expected mass to RLOO’s ( $A_N^{\text{RLOO}} = 2p(1-p)$  exactly, under the  $1/N$ -normalized leave-one-out weights) tends to  $N - 1$ .*

**Interpretation.** The comparison to GRPO is two-sided and exact for the deployed implementation, which normalizes by the **sample** standard deviation ( $s_K = \sqrt{K(N-K)/(N(N-1))}$ , PyTorch’s default), not the population SD:  $A_N^{\text{GRPO}} = 2\sqrt{\frac{N-1}{N}} \frac{1}{N} \mathbb{E}\sqrt{K(N-K)}$  with  $K \sim \text{Bin}(N, p)$  (MC-verified against `estimators.py`). Taking tails,  $\text{MaxRL/GRPO} \rightarrow \sqrt{N}$  as  $p \rightarrow 0$  while  $\text{GRPO/MaxRL} \rightarrow (N-1)/\sqrt{N}$  as  $p \rightarrow 1$  (at  $N=16$ :  $4.0\times$  and  $3.75\times$ ; the symmetric  $\sqrt{N-1}$  holds only under population-SD normalization). MaxRL over-serves the frontier and releases mastered prompts; GRPO silently keeps  $\sim (N-1)/\sqrt{N}\times$  more of its mass on mastered prompts — maintenance a curriculum then removes. Concentrating sampling on the frontier only helps if the estimator extracts signal there, and §6 shows the mastered-tail difference becoming an active failure under a curriculum.

**Proposition 3** (Allocation, static). *Consider the static problem  $\max_{\{N_i\}} \sum_i u_{N_i}(p_i)$  subject to  $\sum_i N_i = B$  and integer  $N_i \geq 0$ , with pass rates  $p_i$  fixed and known and one group per prompt. Because the marginal gain  $u_{N+1}(p) - u_N(p) = p(1-p)^N$  is decreasing in  $N$ , greedy water-filling on this marginal — the probability that the next rollout is a group’s first success — is optimal.*

**Interpretation.** This is a mass-optimal reference allocator for the static, known- $p$  problem only; it does not prove that sampling prompts  $\propto u^\gamma$  is optimal (at fixed group size, expected immediate mass is linear in the sampling distribution, so the unconstrained myopic optimum is a brittle arg-max), nor does it describe the §6.1 oracle, which samples prompts by true- $p$  utility at fixed group size rather than varying  $N_i$ . The practical teacher (§4) keeps fixed group sizes and samples prompts  $\propto u^\gamma$  with a floor, trading myopic optimality for group-parallel

batching and exploration. We measure what the gap costs: a **true-pass-rate utility oracle** (no floor,  $\gamma$ -matched) ties the Thompson teacher in end performance (§6.1).

*Remark 1* (Scope of the surrogate). Coefficient mass counts weighted rollouts, not score-vector geometry. We use it as the sampling utility because (i) it is knowable *before* spending rollouts, (ii) its zeros are exact — at  $p \in \{0, 1\}$  the estimator emits no gradient of any norm — and (iii) on the exact-gradient testbed it *ties* the true first-order expected improvement as a task ranking (Spearman within .005 at every snapshot, and exactly 1.0 against the full exact transfer matrix — one expected group per task, exact  $\Delta p$  on every other task — at three training snapshots; artifacts `results_bridge.json`, `results_transfer_matrix.json`). That tie is a property of *shared-prefix* task structure, where activity and transfer align; on pools without it, activity ranks what a task can emit, not where it helps (§6.5–6.6 measure that gap). What it does not claim to be is the optimal sequential utility, and no state-local functional is: on independent tasks the true final-performance optimum is a planning solution (trajectory-knapsack with a budget-dependent funding cutoff), and which surrogate best approximates it depends on the training objective — a variance-tilted  $(1-p)^{\alpha} u_N$  wins anytime performance while model-predictive replanning wins the final checkpoint, each losing the other’s metric (8–10/10 paired seeds each way; App. A). “Best curriculum utility” is ill-posed until that objective is fixed; we report both currencies throughout. The load-bearing claim is the pair of exact zeros at  $p \in \{0, 1\}$ , which every candidate utility shares.

*Remark 2* (The curriculum changes the training distribution). Sampling prompts from an adaptive  $q_t$  instead of the data distribution  $\rho$  optimizes a reweighted objective  $\mathbb{E}_{x \sim q_t}[J(x)]$ : each group’s gradient is an unbiased MaxRL gradient *for its own prompt*, but the mixture is tilted toward the frontier, without importance correction back to  $\rho$ . This is the standard bargain every curriculum strikes (uniform sampling makes the same choice for  $q_t = \rho$ ), and we treat it as an empirical question, not a free lunch: all evaluations in §6 are on fixed held-out pools under  $\rho$ , so any mismatch cost is charged to the method. The safety results are precisely about when this tilt helps or hurts the  $\rho$ -measured outcome.

One identity, three literatures: learnability curricula (Rutherford et al., 2024; Foster et al., 2025) are the  $N=2$  slice of  $u_N$ ; DAPO-style dynamic sampling and GRESO are its avoidance shadow (cull where  $u \approx 0$ ); HER-style relabeling (Andrychowicz et al., 2017) is its creation complement (manufacture  $K > 0$  where  $u = 0$  identically; §5). DisCO (Li et al., 2025) and SC-SDPO (Liu et al., 2026) derive analogous per-question weights for the GRPO family to diagnose difficulty bias; our identity generalizes that line to success-conditioned estimators and — the step neither takes — uses it to predict how *external* curricula and recyclers interact with each estimator, in coverage terms.

**The three regimes, as a map of the paper (Fig. 3).** The **teacher** (§4) allocates compute within the learnable region — and saturates: a true-pass-rate oracle allocator ties the cheap posterior in end performance (§6.1), so better pass-rate information is not where remaining gains live. **Recycling** (§5) creates signal where sampling the requested task yields none — in our control battery, the one intervention that still adds on top of the oracle. The **estimator** underneath shapes whether either is safe, and its mastered tail is where both failure modes live (§6). One line: *the teacher allocates, hindsight creates, the estimator decides whether either is safe*.

## 4 FRONTIERMAX: THE SCHEDULE

FrontierMax is a standard group-rollout RLVR loop with three modified lines (Alg. 1). A decayed Beta posterior tracks each prompt’s pass rate from observed group outcomes only — never from relabeled successes, which inflate it ( $\hat{p}$  0.81 vs. eval 0.47). Thompson sampling draws  $\tilde{p}$  and samples prompts  $\propto u(\tilde{p})^\gamma$  with a uniform floor. The knobs (decay 0.7, floor 0.1,  $\gamma \in \{1, 4\}$ ) carry validated defaults (Table 3); what is *derived* rather than tuned is the difficulty band itself — location, width, and  $N$ -scaling — which is exactly where competing curricula spend their hyperparameters.

## 5 FAILURE RECYCLING

*Remark 3* (Hindsight update, characterized). Relabeling a dead group to the sub-goal  $g'$  its trajectories actually achieved, rewriting the goal conditioning, and applying the same success-conditioned weights yields a *verified weighted-likelihood update for the achieved task*: every relabeled example is a true






Figure 2: **One FrontierMax step.** A Thompson-sampled posterior allocates rollouts by the derived utility; live groups train with unmodified MaxRL weights; dead groups are relabeled to the sub-goal they actually achieved (conditioning rewritten) and train as verified successes. The posterior observes requested-task outcomes only.

**Algorithm 1** FrontierMax: one step. Three lines differ from a standard group-rollout loop (marked  $\triangleright$ ).

- 1:  $\triangleright$  **sample** prompts  $x_1..x_B \propto u_N(\tilde{p}_x)^\gamma$  mixed with a uniform floor,  $\tilde{p}_x \sim \text{Beta}(\alpha_x, \beta_x)$  (else: uniform)
- 2: roll out  $N$  completions per prompt; verify;  $K_x \leftarrow$  successes
- 3: update posterior  $(\alpha_x, \beta_x)$  from  $(K_x, N)$  for the requested task only
- 4: live groups ( $1 \leq K_x$ ): MaxRL advantages  $w_i = r_i/K - 1/N$ , unchanged
- 5:  $\triangleright$  **dead groups** ( $K_x=0$ ): relabel each parseable failure to its own achieved sub-goal  $g'$  (verified; conditioning rewritten); unparseable rollouts stay failures, so the group retains contrast ( $K < N$ )
- 6:  $\triangleright$  admit the relabel only if  $\hat{p}_{g'} \leq \text{gate\_max\_p}$  (the gate: refuse the mastered tail)
- 7: one policy-gradient step on live + admitted relabeled groups

success of  $g'$  under the environment’s verifier. It is generally an *adaptively selected, off-policy auxiliary update*, not an unbiased gradient of  $g'$ ’s on-policy objective:  $g'$  is selected from the same group being scored, conditioning on the requested task’s failure correlates the rollouts, the conditioning rewrite changes the law under which the log-probability is evaluated, no importance ratio corrects for it, and the finite- $N$  weights target  $J_{N-1}$  (Lemma 1) rather than exact ML. It equals the fresh-rollout  $J_{N-1}$  gradient for  $g'$  only under a distribution-matching condition — the selected, rewritten group must have the same joint law as an i.i.d. fresh rollout group for  $g'$ . This specializes known relabeling-bias structure — the induced-distribution view of Eysenbach et al. (2020) and the conditional-law gap in Ghosh et al. (2021) — to success-conditioned RLVR weights; we state it as a remark, not a result, and note that the sharpening phenomenon of §6.8 runs through the *marginal shift* those works also identify, not this conditional matching.

**Interpretation.** “Exactness” is a property to measure per environment, not a blanket guarantee. On our validation testbed the laws match by construction, so cosine probes there are consistency checks, not evidence (several null relabeling schemes produce identical cosines); the evidence that the relabel direction carries information is behavioral — the placebo decomposition below. We measured the two gaps the remark names where they are measurable (exact chains, 5 seeds; artifact `results.proposal.shift.json`): the applied relabel update’s cosine against the exact fresh-destination gradient direction is 0.93, the relabeled group’s success rate exceeds the true fresh destination pass rate by only +.003 (selection guarantees  $K' \geq 1$ , so the bias exists but is small here), and cross-fitting the destination choice (select on half the group, update with the other half) removes the adaptive reuse at a real cost: alignment drops to 0.82 and AUC falls .883  $\rightarrow$  .829, because the selected successes leave the update half. The off-policy caveat is real; its sign is environment-specific.

Two contracts keep practice on the proof’s side: (1) **exactness** — a relabeled success must be a true success of the relabeled task under the environment’s own verifier (never an LLM judge); (2)




Figure 3: **The three regimes as a map of the method.** The teacher allocates compute within the learnable region; recycling creates signal where sampling the requested task yields none; the gate refuses relabels that land in the mastered tail; the estimator underneath decides whether any of it is safe.

**conditioning** — goal-embedding trajectories must have their conditioning rewritten to the achieved goal. Skipping (2) turns the exact gradient into noise: on our gridworld, hindsight without the rewrite scores below teacher-only ( $0.600 < 0.658$ ); with it, it leads ( $0.703$ ). Among the interventions studied here, recycling is the one that *creates* signal for an all-fail requested task (adding the achieved sub-goals to the pool by hand would too — what recycling buys is discovering them, §6.2), and we decompose its headline gain rather than let it inflate: of the  $+0.22$  AUC on fixed task pools, a placebo replaying *live* gradients on dead-group slots captures 83% — most of the benefit is extra gradient dose on otherwise-wasted slots — while the relabel *direction* carries  $+0.037$  beyond the strongest placebo, is the tightest arm in the battery, and is exactness-dependent (random-target relabels lose to dose-matched replay 5/5). Both numbers ship together. The value is regime-dependent in a way we can state precisely:  $+0.22$  AUC on fixed task pools,  $+0.01$  on one-shot task streams — the gain is proportional to how much a relabeled skill can compound, and task-set fixedness is the single most important regime variable we found.

## 6 EXPERIMENTS: AN ESCALATING LADDER

Each rung asks one question of the three-regime structure, at increasing scale and decreasing experimental control; Table 1 is the map. Statistical contract, stated once and used throughout: the independent unit for any training-method claim is an independently trained seed block; runs sharing a seed, warmstart, or checkpoint are paired or correlated observations, never additional replicates; levels,  $k$  values, metrics, and checkpoints from one trained model are repeated measurements of it; repeated-evaluation spread measures evaluation noise, not training-seed uncertainty, and every  $\pm$  names its source at point of use ( $\pm$  with no qualifier is standard deviation across training seeds). “Coverage” is  $\text{pass}@k$  with exactly  $k$  samples drawn per prompt and averaged over the eval pool, at the  $k$  named per suite (8 on the maze, 16 on Countdown, 4 on GSM8K);  $\text{mean}@k$  is the mean success fraction of the same  $k$  samples; AUC is the mean of the eval curve over training at matched wall-clock or episode budget, and final is the last checkpoint. Where run pools are heterogeneous we say so and report signs descriptively rather than attaching  $p$ -values (§6.3–6.4). Pre-registered predictions were committed before the deciding runs; negative results appear in place, one sentence each, with full postmortems in the repository’s audit trail.

**6.1 Does allocation saturate? (exact-gradient skill chains; CPU, 5 seeds).** Uniform  $0.650 \rightarrow$  teacher  $0.728 \rightarrow$  full stack  $0.890 \pm 0.002$  AUC. An eight-arm control battery (App. A) orders into a *gradient-information ladder*, each rung a 5/5 paired-seed win: zero-information dose ( $0.798\text{--}0.800$ )  $<$  replayed correct direction ( $0.832$ )  $<$  true relabels ( $0.869$ )  $<$  perfect allocation (oracle  $0.8885 \approx$  full stack  $0.8895$ )  $<$  allocation+creation ( $0.8935$ ). Two decompositions matter downstream. *Allocation*

Table 1: The ladder at a glance: every rung asks one question of the three-regime structure; scale rises, control falls.

| rung | question                      | setting                   | seeds   | outcome                       |
|------|-------------------------------|---------------------------|---------|-------------------------------|
| 6.1  | does allocation saturate?     | skill chain (exact grads) | 5       | yes: oracle ties posterior    |
| 6.2  | anything to sample toward?    | frontier-heavy pool       | 5       | no: only creation scores      |
| 6.3  | main effect, real grads?      | maze transformer, 1.26M   | 9 runs  | 9/9 signs (exploratory)       |
| 6.4  | where/when does coverage go?  | maze, level-resolved      | 22 runs | easy band, immediately        |
| 6.5  | when does allocation not pay? | IsaacLab Anymal-C         | 1       | learnable-everywhere: null    |
| 6.6  | beyond task spread?           | MountainCar/CartPole      | 3       | spread=existence; tea.=speed  |
| 6.7  | LLM scale?                    | GSM8K 2×2, 360M           | 1–2     | sign landed; replication open |
| 6.8  | what does recycling cost?     | Countdown v2              | 3       | mean up, coverage down        |
| 6.9  | does the gate recover it?     | Countdown v2              | 3       | one useful trade point        |




Figure 4: **The ladder.** (a) Exact-gradient chains: a no-floor,  $\gamma$ -matched true-pass-rate oracle matches the full stack within the observed five-seed variation — allocation saturates — and recycling still adds on top of the oracle. (b) Frontier-heavy regime: allocation scores exactly zero; recycling solves it, and does so equally under uniform sampling (.931 vs. .928) — creation, not allocation, is the live channel. (c) Maze at matched wall-clock: 3/3 seeds on both meters. (d) GSM8K 2×2, registered run, error bars = measured same-model eval noise (SD .0094): GRPO+teacher is the only cell that regresses (pre-registered P-G2; a replication seed did not reproduce it — §6.7); the teacher’s MaxRL effect is a null at this budget (P-G1).

*saturates:* a no-floor,  $\gamma$ -matched true-pass-rate oracle ( $0.8885 \pm .0014$ ) sits within the five-seed variation of the full stack ( $0.8895 \pm .0023$ ) — similar within observed noise, not tested for equivalence — so perfect difficulty information is worth little over the cheap posterior. *Creation decomposes:* a replay placebo captures 83% of recycling’s  $+0.22$ ; the relabel *direction* carries  $+0.037$  beyond it and is the most reliable arm in the battery, while random-target relabels lose to dose-matched replay 5/5. The utility form is load-bearing through its  $N$ : the  $N$ -blind learnability slice  $p(1-p)$  forfeits the entire allocation gain ( $-.007$  vs. uniform) where  $u_N$  carries  $+0.078$  (5/5). All arms converge within .027 of one another — on fixed pools these interventions buy *time-to-ceiling*, not the ceiling.

**Takeaway.** A monotone information ladder: dose < direction < allocation-ceiling < +creation. Allocation saturates; creation adds on top; every intervention buys speed, not asymptote.

**6.2 Can any sampler act when nothing is learnable? (frontier-heavy regime, max pool pass rate  $10^{-5}$ ; 5 seeds, artifact committed).** Uniform, DAPO, and the plain teacher score **exactly 0.00**; adding recycling reaches **0.98 final (0.93 AUC)** — relabeling invents the curriculum below the pool, ignites within  $\sim 400$  groups, then goes silent. The attribution control pins down which channel did it: *uniform+recycling* scores the same ( $0.931 \pm .002$  vs. *teacher+recycling*  $0.928 \pm .003$  AUC) — in this regime the entire effect is signal creation; allocation contributes nothing because there is nothing



learnable to allocate toward. Equivalently: a designer who hand-added the achieved sub-goals to the pool would also escape the zero — what recycling buys is *discovering* that sub-pool from the policy’s own failures, with no task engineering. One loss-level alternative remains: the MaxRL paper’s own Eq. (10) keeps the control-variate term  $-\frac{1}{N} \sum_i S_i$  on all-fail groups, so that estimator is *not* silent at  $K=0$  ( $\mathbb{E}[\hat{g} \mid K=0] = \nabla p/(1-p)$ , MC-verified). We ran it (5 seeds, matched budget; artifact `results_fullcv_baseline.json`): on this pool the full-CV variant spends 100% of its update norm on all-fail groups and still scores 0.000 in every seed, with or without the teacher — the all-fail term pushes down every sampled action symmetrically and cannot ignite a pool this deep, while recycling under either variant reaches 0.98. On the balanced pool the two variants tie (paired  $\Delta$ AUC straddles zero, 5 seeds).

**Takeaway.** Where nothing is learnable at budget, sampling policy is irrelevant — uniform and teacher tie at zero without recycling and tie at 0.93 with it. Even the estimator family’s nonzero all-fail update (full control variate) stays at zero here: creation is the only channel we tested that ignites a dead pool, and its value is discovery, not allocation.

**6.3 Does the estimator main effect survive real gradients? (maze transformer at matched wall-clock; 1.26M params, 13 levels,  $\sim 30$  runs, 3 seeds).** Champion (teacher + dense recycling)  $0.252 \pm 0.005$  final /  $0.229 \pm 0.009$  AUC vs. uniform  $0.230 \pm 0.015$  /  $0.211 \pm 0.011$ ; paired deltas positive in 3/3 seeds on both meters (final and AUC — six correlated contrasts, not six seeds); 22–35% more optimization steps per GPU-hour. Attribution within that gain, stated plainly: matched *wall-clock* lets curriculum arms take 1.2–1.4 $\times$  more optimization steps (dead groups skip the backward pass), and step-matched truncation shows the extra steps are most of the pure teacher’s effect (step-matched teacher-only  $\Delta$ AUC +0.006, within seed noise) while teacher+recycling survives step-matching (positive in all three paired seed blocks; a paired  $t$  at  $n=3$  gives  $p \approx .04$ , which we report as descriptive given its normality sensitivity — an exact sign-flip test cannot go below .25 at three blocks) — on the maze the teacher’s measured value is *throughput* (buying more steps per GPU-hour), recycling’s is signal. The maze teacher uses the legacy frontier form; a pre-registered 3-seed rerun with the exact  $u_N$  measures the two within noise *on the maze itself* (AUC  $.2138 \pm .0059$  vs.  $.2135 \pm .0104$ ; artifact `un_uniform_verdicts.json`), closing the gap between the deployed form and the derived one. The same sweep refutes a CPU-derived within-band refinement (a  $(1-p)$ -tilted utility that won 10/10 on the testbed loses 0/3 on the maze on every meter) — within-band utility choices are testbed-sensitive, which is why the theory claims only the band’s boundaries. Safety, stated as what the design supports: an estimator-*associated* coverage difference — all five MaxRL-labeled runs grow `pass@8` (e.g.  $0.316 \rightarrow 0.348$  under the teacher) and all four GRPO-labeled runs lose it ( $0.308 \rightarrow 0.271$  uniform;  $0.332 \rightarrow 0.269$  with the teacher). We state this as what it is: perfect sign separation across nine runs whose pool mixes sampler configurations (the MaxRL five span frontier and uniform teachers; the GRPO four span uniform and one frontier) and reuses seed blocks and warmstarts, so the runs are not exchangeable and a label-permutation  $p$ -value would answer a randomization this design did not perform — we therefore report zero exceptions descriptively and do not attach one. Two sensitivity reads support the same direction without curing the design (artifact `permutation_reanalysis.json`): restricting permutations within matched sampler arms preserves the perfect separation, and all three shared-warmstart pairs are individually positive; the run composition (5v4 because one uniform-MaxRL seed predates `pass@k` logging) is tabulated in the same artifact. The confirmatory design this motivates is a balanced  $\{\text{MaxRL, GRPO}\} \times \{\text{uniform, teacher}\}$  factorial over at least six independent seed blocks with identical warmstarts and one prespecified endpoint — fewer than six blocks cannot reach two-sided  $p < .05$  even with unanimous signs. At the exact-gradient rung, where the design is cheap, we already ran its schedule-matched core: freeze each seed’s *realized* prompt sequence (one recorded teacher trajectory, one uniform), then replay the identical sequence under each estimator from the same init — the estimator ordering on exact summed `pass@8` survives with the schedule held fixed (MaxRL over GRPO in 10/10 seed $\times$ schedule pairs; artifact `results_schedule_matched.json`), with the caveat that on that learnable-everywhere pool GRPO still gains coverage — the neural-scale *loss* is what the balanced maze factorial must test. Whether the teacher *amplifies* GRPO’s decay is a factorial (interaction) question this suite is not powered for; the GSM8K  $2 \times 2$  below tests exactly that cell, pre-registered. Inference currency (one convention, every level: smallest  $k$  reaching absolute coverage 0.25; single checkpoint pair): our checkpoint needs  $11 \times$  fewer samples at level 5, ties GRPO at levels 2–3, and needs  $3 \times$  more at level 4 — the frontier is where the efficiency gain lives.




Figure 5: **Where and when GRPO’s coverage goes (maze, level-resolved, matched wall-clock).** (a)  $\Delta \text{pass@8}$  (final – warmstart) by difficulty band: GRPO’s loss lives on the easy/mid band (all 4 runs; dots are runs) while its pass@1 there rises — sharpening; MaxRL holds the easy band and gains on the frontier band (18 runs; a heterogeneous run inventory — descriptive, no  $p$ -value). (b) Diversity premium  $\text{pass@8} - \text{pass@1}$  vs. wall-clock (mean and min–max across runs): GRPO’s erosion begins within the first  $\sim 25$ –50 steps, before any level saturates; MaxRL stabilizes after warmup.

The summed-coverage view agrees without a target choice: GRPO ahead at pass@1, curves crossing at  $k \approx 4$ , MaxRL ahead by half a level at  $k=64$  (App. A) — choose the estimator by how many samples you will draw at inference.

**6.4 Where and when does GRPO’s coverage go? (level-resolved).** Decomposing the main effect by difficulty band turns it into a mechanism (Fig. 5). *Where:* GRPO’s pass@8 loss is concentrated almost entirely on the easy/mid band L1–3 ( $-.224$  mean over its four runs, every run negative) while its L1–3 pass@1 *rises* ( $+.02$  to  $+.14$ , every run) — it funds reliability by liquidating easy-band coverage; MaxRL holds L1–3 flat ( $+.005 \pm .058$  over 18 runs) and concentrates its gain on the frontier band L4–6 ( $+.126 \pm .063$  vs. GRPO’s  $+.057$ ). The 18-run MaxRL pool is a heterogeneous run inventory (teacher variants, hindsight doses, a wide model, repeated seeds), so we report the band asymmetry descriptively — every GRPO run negative on L1–3, the MaxRL spread centered at zero there — and attach no permutation  $p$ -value; the 22 runs are not exchangeable replicates of one estimator contrast. This is the observable consequence of the mass curves in Fig. 1: the  $\sqrt{K(N-K)}$  weighting keeps spending updates on mastered prompts (sharpening them), where  $u_N$  has already released them. *When:* the divergence is immediate, not saturation-triggered — GRPO incurs 59–73% of its total diversity-premium loss ( $\text{pass@8} - \text{pass@1}$ :  $.123 \rightarrow .042$ ) within the first quarter of the budget, with the first easy-band pass@8 drop visible at steps 0–50 when L1 pass@1 is still only  $.62$ – $.75$ ; MaxRL’s premium stabilizes after warmup ( $.122 \rightarrow .101$ ). A dose-response note from the same logs (descriptive, same caveat as above): sparse relabeling (1/dead group) uniquely preserves the diversity premium relative to none and dense, while dense banks the same coverage into the best final pass@1 — dose sets *where* the coverage dividend is spent, not whether coverage is lost.

**Takeaway.** The gains survive real gradients, and the coverage main effect now has a located mechanism: GRPO liquidates easy-band coverage for reliability from the first updates; MaxRL reallocates the same budget to the frontier. Coverage is the currency that sees it.

**6.5 When does allocation *not* pay? (IsaacLab legged locomotion, pre-registered null).** A 5-arm Anymal-C rough-terrain pilot (one seed, run by an independent team on a co-designed protocol; App. A) lands its pre-registered null exactly: on a terrain grid where every level is learnable at budget, the stock greedy curriculum beats the frontier teacher outright (mean pass  $.372$  vs.  $.278$ ). Read with the maze step-matched analysis and Countdown, the consistent regime pattern — at one seed here, so a pilot-level observation rather than proof — is: **allocation pays where unlearnable-at-budget regions exist to avoid**; on learnable-everywhere pools, breadth beat focus in every case we measured.




Figure 6: **GSM8K 2×2 (pre-registered; registered run shown; error bars = measured same-model eval noise, SD .0094).** GRPO+teacher is the only cell that regresses in its second half (P-G2 on this run; the load-bearing statistic is the endpoint gap vs. plain GRPO, 2–4× the repeated-eval noise — a descriptive signal-to-noise read, not a test; the second-half dip alone is within noise; a replication seed with weaker measured steering climbed instead — §6.7); under MaxRL the identical teacher trains stably to its best value but does not beat uniform at this budget (P-G1: null).

**6.6 What does a curriculum add beyond task spread? (Gymnasium control — binary-success RLVR on classic control; official-task currency).** All arms train on *binary task success* (MountainCar: reach the flag; CartPole: survive 400) — the classic-control analogue of a verifier reward. The **target-only** arm is then a provable zero-update freeze, not a slow learner: every group is all-fail, every group-relative estimator emits exactly zero weights on constant rewards, and telemetry confirms zero parameter updates over the full budget (random-policy success  $< 3 \times 10^{-6}$ ; no budget ignites it). The gated full stack under the identical reward reaches **1.000±0.000 on both official tasks** (3 seeds, Fig. 8). A dense-reward control confirms the zero is a property of sparse verifier-style reward (REINFORCE on CartPole’s native reward solves it; App. A). Attribution at a common horizon: uniform-over-bins task spread supplies most of the effect (existence of a gradient), recycling-without-allocation plateaus below ceiling, and the teacher’s remaining contribution is speed and stability (3× faster to hard-task 0.9, 2.5–6× smaller cross-seed variance). A methods note we keep because it generalizes: our own first version used per-bin task-partitioned policy parameters and reproduced the known failure (0.000 in every arm) — curricula operate through shared parameters, or not at all.

**Takeaway.** Curricula operate through shared parameters, or not at all.

**6.7 Does it transfer to LLM scale? (GSM8K 2×2, pre-registered; SmolLM2-360M,  $N=16$ , one A10G).** Predictions committed before any cell finished. **P-G2 landed on the registered run:** GRPO+teacher was the only cell that regressed in its second half, and the deficit against plain GRPO at the endpoint is 2–4× the measured same-model eval noise (mean@4 −.027; pass@4 −.048; the second-half dip itself, .096 → .093, is within the ±.009 noise floor and carries no weight on its own) while plain GRPO climbed monotonically (.078 → .105 → .120). One uncertainty caveat governs every GSM8K number: the ±.009/±.017 bars are repeated-eval noise of a fixed checkpoint — they bound evaluation variability, not training-seed variability — so deficit-to-noise ratios are descriptive signal-to-noise reads, not hypothesis tests; training-seed uncertainty at  $n=2$  is what the queued multi-seed cell measures. Steering telemetry, read against its own controls: the teacher’s sampled-dead fraction *minimum* reaches .48 vs. uniform’s .52, but the run means are indistinguishable (.66–.67 in all cells) — the delivered treatment was weak, which is the posterior-starvation mechanism below, and it bounds how much of the GRPO gap the curriculum can own. **A replication seed does not reproduce the regression shape:** seed 2 of the same cell climbs monotonically (.073 → .095 → .118), and its steering telemetry shows the treatment was weaker (min dead-sampled fraction .53, mean .66 ≈ population — a near-uniform teacher on average). The completed *within-seed* pair (grpo-uniform




Figure 7: **Recycling-induced sharpening and its derived mitigation (Countdown, tier 1).** (a) Three seeds, both axes: recycling (B2) buys mean@16 with pass@16 coverage. On this tier the moderate gate (B3) gives back most of the mean without moving coverage — tier-1 destinations saturate early, so admitted relabels are few — while full-strength gating (corrected decay, 1 seed) returns coverage *above* baseline at the mean’s full price. The settings are consistent with a mean-vs-coverage trade (not yet a monotone dial: the strong point is one seed on different code), and the useful middle lives on the frontier tier, where the same moderate setting restores coverage to baseline ( $.279 \pm .019$  vs.  $.274$ ) keeping  $\sim 60\%$  of the mean gain ( $\$6.8$ ). (b) The mechanism in-run: the gate’s rejection rate rises monotonically as the model’s reachable set saturates its relabel destinations, while the ungated dose rides its cap. (c) Gated recycling preserves *more* policy entropy than no recycling — admitted relabels are a spread-preserving gradient.

seed 2:  $.081 \rightarrow .085 \rightarrow .125$ , final pass@4  $.229$ ; artifact `grpo_uniform_seed2.json`) sharpens the read: the endpoint teacher deficit under GRPO has the same *sign* in both seeds on both meters (mean@4 and pass@4 are computed from the same generations and strongly correlated, so this is two seed-level confirmations read on two coupled meters, not four; seed 2:  $-.007$  mean@4, within noise,  $-.026$  pass@4,  $\sim 1.5 \times$  its floor) with magnitude tracking steering intensity — a dose effect, not a sign flip. The pre-registered *second-half-regression* shape remains 1 of 2 seeds: suggestive, not established. MaxRL+teacher trains stably to its best value ( $.066 \rightarrow .102$ ); the uniform-MaxRL cell finishes at  $.108$ , making P-G1 (a teacher gain for MaxRL) a *null at this budget* — predicted in advance by the posterior-starvation analysis below. A pass@ $k$  sweep on the registered checkpoints: the teacher’s deficit under GRPO *widens monotonically with  $k$*  ( $-.008$  at  $k=1$  to  $-.024$  at  $k=16$ ) — curriculum damage lives in coverage, as the maze predicted, though each individual delta is within single-seed noise (harness caveats in App. A). One design confound is disclosed rather than dismissed: the teacher samples with replacement, so curriculum cells revisit prompts more — a repetition effect is an alternative reading that a uniform-with-replacement control cell separates; it is running with the steering-controlled replication. The mechanism behind the weak treatment: the teacher’s posterior *learns* real difficulty from  $\sim 1$  visit per prompt ( $\rho = -0.17$  against independent 7B-model solve-rate annotations,  $p \approx 10^{-17}$ ; cross-run posterior agreement  $3 \times$  its agreement with the external key — it learns *this model’s* landscape) but its allocation is *posterior-starved*: 3,200 draws over 7,473 prompts leave the estimated utility landscape nearly flat, so Thompson sampling stays near-uniform (full posterior diagnostics in App. A). Small model; the pre-registered outcomes were orderings and signs on the registered run.

**Takeaway.** The maze sign pattern stands on nine heterogeneous runs; the LLM-scale interaction is 1-of-2-seeds and tracks steering intensity — the decisive experiment is a steering-controlled multi-seed cell, queued.

**6.8 What does recycling cost? (Countdown  $2 \times 2$ , pre-registered): two nulls and a new phenomenon.** The first exact-verifier hindsight experiment in RLVR ran in production (12 relabeled groups/step, 72–89% yield). Both positive predictions returned null, with an exact diagnosis: the

uniform baseline reached 94/76/94% of the random-guesser bound implied by the pool’s simplified construction — no headroom for allocation or creation to matter. What the run found instead is new: **recycling-induced sharpening** — hindsight arms improved mean@16 while *losing* pass@16 coverage (tier-1: .571 vs. .645 for the no-relabel baseline). Recycled off-target successes concentrate the policy on reachable outputs at the expense of exploration breadth: the creation channel has its own safety trade-off, mirroring §6.3’s weight-induced collapse — data-induced rather than weight-induced. The underlying concentration pathology is known in goal-conditioned RL (Skew-Fit’s motivating failure mode; GCSL’s marginal shift); what is unreported, to our knowledge, is its measurement in RLVR — no prior work quantifies the pass@ $k$  cost of hindsight relabeling in a verifier loop. Our algebra flags the mechanism: the achieved-goal marginal overweights destinations the policy already reaches, so relabel updates concentrate toward the mastered tail, where  $u(p)$  says fresh training signal vanishes — the softmax pays for those updates out of the exploration tail. (One selected success does not by itself imply the destination’s fresh-rollout pass rate is high; that is what the gate’s independent destination estimate  $\hat{p}_g$  checks.) *Takeaway: recycling needs an admission rule, and the utility’s high- $p$  zero motivates one — a destination-saturation threshold on the estimated pass rate  $\hat{p}_g$  (the implementation thresholds  $\hat{p}$  directly; it does not evaluate  $u_N$ ).*

**6.9 Does the gate recover it? (Countdown v2 pool, 3 seeds, pre-registered).** We rebuilt the pool with real structure (random operand permutation, parenthesization, exact division; a headroom gate excluded guesser-saturated tiers) and ran three arms  $\times$  three seeds: no recycling (B1), ungated recycling (B2), gated recycling (B3). Setup note with a theory reading: recycling has a *syntactic prerequisite* — the pre-SFT base model produces zero verifier-valid failures (relabel yield 0.000,  $n=384$ ), so there is nothing to recycle until SFT installs the output format; post-SFT the model is a near-independent guesser whose pass@16 sits at a tier-invariant  $\sim 75\%$  of the independence bound  $1 - (1 - \text{pass}@1)^{16}$  (ratios .744/.783/.755 across tiers) — the baseline our headroom gate measures against. **Sharpening replicates:** at tier 1, recycling buys mean@16  $+ .046$  (.278 $\pm$ .054  $\rightarrow$  .324 $\pm$ .012) and pays  $-.049$  of pass@16 (.541 $\pm$ .020  $\rightarrow$  .492 $\pm$ .011). The coverage loss clears both arms’ seed spreads ( $> 2\times$ ); the mean gain clears the treated arm’s tight spread but sits within the noisier baseline’s — we report it as directionally consistent across seeds rather than individually significant. The frontier tier moves the same direction. **The timing is dose-then-damage:** recycling’s frontier-tier mean gain is installed by step 30 —  $+ .014\pm .008$  paired, all seeds positive, at only  $\sim 28\%$  of the eventual relabel dose — while the coverage cost only materializes by step 45 ( $-.048\pm .019$  paired, *all seeds negative*, after  $\sim 65\%$  of dose): the benefit precedes the damage, so early-stopping on mean-accuracy validation would bank the gain and never see the bill. The damage is also visible on the training side: the ungated arm ends with more unsolvable-at-16 frontier training prompts than baseline in every seed; the gate’s entropy retention replicates across seeds while ungated recycling makes entropy seed-unstable (details in App. A). **The gate works where the algebra says it should:** at the frontier tier (destinations not yet saturated) the gated arm returns coverage to baseline (.279 $\pm$ .019 vs. .274 $\pm$ .015) while keeping  $\sim 60\%$  of the mean gain; at tier 1 (destinations saturate early, so the gate blocks nearly all relabels) it neutralizes recycling almost entirely. Implementation disclosure: a decay-math bug made all three-seed gated runs gate *weaker* than designed (effective destination threshold  $\hat{p} \approx 0.64$  where the design implies 0.92), so the reported operating point is an under-gated one; the corrected-decay rerun (1 seed) gates  $\sim 3\times$  harder and pushes tier-1 coverage *above* baseline at the mean’s full price (Fig. 7a). The three settings are *consistent* with a mean-vs-coverage trade — no gate: full mean gain, coverage lost; under-gated:  $\sim 60\%$  of the gain, frontier coverage restored; full-strength: gain erased, coverage above baseline — but they do not establish a monotone dose-response: the strong point is one seed, gate setting is confounded with the decay-code version, and at tier 1 the moderate gate moved *both* point estimates below the ungated arm. We claim the moderate frontier-tier setting as one useful operating point and defer the dial claim to the pre-registered fixed-decay sweep. In-run telemetry shows the mechanism (Fig. 7b–c): rejection rises .12  $\rightarrow$  .85 as the reachable set saturates, and the gated arm ends with *more* policy entropy than the no-recycling baseline. A pre-registered gate\_max\_p dose sweep at fixed decay is the standing follow-up.

**Takeaway.** Sharpening is a reproducible property of exact-verifier recycling, and the derived gate buys back frontier-tier coverage at a measured mean price — one validated operating point, with the monotone-dial question deferred to a fixed-code sweep.



## 7 RELATED WORK

**The coverage tension.** RLVR itself trades coverage for precision:  $\text{pass}@1$  rises while  $\text{pass}@k$  at large  $k$  falls below the base model (Yue et al., 2025), with entropy collapse as the mechanistic signature (Cui et al., 2025) and overtraining on saturated problems as a locus (Yuan et al., 2026). Our contribution to this line is the *interaction*: whether external curricula and recyclers compound or cancel the estimator’s implicit weighting, and what that costs in coverage. Notably, recent work frames curricula as the *cure* for coverage loss (Cai et al., 2026); our GRPO arms show the cure is estimator-conditional — those studies never varied the estimator. The estimator–difficulty interaction itself is a heating 2026 theme — SC-SDPO (Liu et al., 2026) shows GRPO’s normalization implicitly encodes a difficulty weighting, and A-GRAE (Yu et al., 2026) redesigns the advantage around difficulty focus — but no prior work runs *identical* curricula under two estimators or measures the resulting  $\text{pass}@k$  sign flip.

**Curricula for RLVR.** Difficulty-adaptive sampling is bucket- or scalar-level and heuristic-scored: DUMP (UCB over buckets), AdaRFT (scalar controller over precomputed labels, TMLR 2026), SEC (category bandit), threshold curricula. LILO (NeurIPS 2025) is the closest per-prompt method — rejection sampling by  $\hat{p}(1 - \hat{p})$ , which Prop. 1’s corollary identifies as the  $N=2$  slice of our derived utility. Reinforce-Ada (Xiong et al., 2025) — the adaptive sampler MaxRL itself points to — adapts the rollout count *within* a prompt until a signal criterion is met, holding the prompt distribution fixed; our teacher is the complementary move, reallocating *across* prompts at fixed group size, and Prop. 3 is the idealized form of both (water-filling spends its next rollout wherever  $p(1 - p)^N$  is largest, across or within). DAPO’s dynamic sampling and GRESO cull dead prompts — avoidance without allocation or creation. PKPO and Pass@k-training modify the objective toward  $\text{pass}@k$  but keep uniform sampling. **Goal selection in HER.** Selecting *which* goals to train on by difficulty or reachability is a decade-old idea in the relabeling literature: intermediate-difficulty goal generation (Florensa et al., 2018), value-based hindsight goal selection (HGG; Ren et al., 2019), curriculum-guided HER’s proximity–diversity admission of relabel goals (Fang et al., 2019), and Skew-Fit’s  $\alpha$  as a concentration-vs-coverage dial (Pong et al., 2020). Our gate belongs to this mechanism class; what is new is not admission-by-difficulty but its *derivation* — the threshold is the mastered-tail zero of the same estimator functional  $u_N$  that schedules sampling, so teacher and gate are one object with no separately tuned goal-selection heuristic — and the coverage accounting it is validated against.

**Hindsight and self-training for LLMs.** Exact-verifier relabeling has precedent outside RLVR loops — CodeIt (Butt et al., 2024) (ARC program synthesis, relabel + prioritized replay) and proof-tree relabeling (Poesia et al., 2024) — and LLM-judged variants relabel agent trajectories (AgentHER, HSL). Concurrent 2026 work brings relabeling-adjacent mechanisms into RLVR-style training: LfH (Xu et al., 2026) (a VLM proposes one shared hindsight instruction per failed group for VLA post-training; the relabeler is a judge, their stated limitation) and ZipRL’s hindsight response replay (Hu et al., 2026) (densifying GRPO signal for multi-turn agents). None of these measures  $\text{pass}@k$  under relabeling — precisely the currency where we find recycling’s hidden cost. The nearest phenomenon report is in same-problem rejection finetuning: Strozzi (2026) document a STaR-trained model winning  $\text{pass}@8$  while the base model overtakes at  $\text{pass}@256$  — concentration over reach — for SFT on a model’s own verified outputs. Recycling-induced sharpening is the cross-task, on-policy-RL sibling, and our gate is its estimator-derived admission rule; gating *vanilla* RLVR updates by saturation signals has independent precedent (Yuan et al., 2026’s Bayesian boundary gating), not applied to relabel destinations. So the precise claims we retain: no prior work (i) derives the sampling utility from the estimator’s own expected signal, (ii) couples it with success-conditioned weights it provably matches, or (iii) *measures the coverage cost of hindsight relabeling in RLVR* and gates relabel admission by destination saturation.

## 8 LIMITATIONS AND NEGATIVE RESULTS

**Scale.** The LLM-scale teacher×estimator interaction is 1-of-2 seeds and tracks steering intensity: the replication seed of the GRPO+teacher cell did not reproduce the registered regression, and its telemetry shows the treatment was barely delivered (§6.7). The GSM8K teacher is also coverage-neutral-to-negative under *both* estimators at this budget, so the contrast there concerns the sign of the mean-accuracy effect, not a coverage rescue. The maze coverage divergence (nine runs, zero sign exceptions, but a heterogeneous pool — §6.3) is the strongest evidence this paper has for the safety

claim, and it is exploratory; the balanced maze factorial and a steering-controlled multi-seed LLM cell are the decisive next experiments.

**Scope of the surrogate.** Coefficient mass is not gradient norm, SNR, or expected improvement (Remark 1) — it omits score-vector geometry, sequence length, token masking, optimizer state, and clipping — and §6.5 exhibits the regime where reallocating it buys nothing (learnable-everywhere pools). Recycled-gradient exactness is a property of the environment’s relabel map; noisier verifiers will land between our fixed-pool and one-shot endpoints (+0.22 vs. +0.01 AUC).

**The coverage main effect is pool-conditional.** A pre-registered fifth suite (Jugs water-measuring puzzles; preregistration, per-cell raw endpoints, verdicts, and postmortem vendored under `jugs_11m/results/` with execution-commit provenance) returned all-null: on a pool whose learnable band is a single stratum of 1–2-move template solutions, *plain MaxRL* also collapses — entropy  $1.36 \rightarrow 0.07$ – $0.21$  in 60 steps, frontier-tier  $\text{pass@16}$   $0.15 \rightarrow 0$ , every arm converging to the same deterministic policy — and recycling accelerates the collapse rather than rescuing it. When  $N$  rollouts of a near-trivial plan are near-identical, the i.i.d.-rollout diversity the group contrast requires is absent, and success-conditioned weights liquidate frontier coverage through shared parameters just as GRPO’s do elsewhere. The estimator main effect of §6.3–6.4 is therefore a claim about pools with a graded band at the deployed  $N$ , not an unconditional safety guarantee; per-task rollout-set diversity at initialization is the design gate we now check first. The same suite also exposed a granularity bug in the gate (keying destinations by value collides tasks when achieved values are few and shared —  $\leq 20$  integers in Jugs vs. hundreds in Countdown); the corrected gate keys on the relabeled task.

**What did not transfer.** A  $4\times$  duration run refuted our own duration hypothesis (the deep-frontier stall is a per-step-competence ceiling; depth needs capacity, not schedule);  $\gamma$ -concentration did not transfer from chains to the maze, as the compounding model predicted; adaptive truncation order underperformed fixed  $T$ ; and feeding relabeled successes to the teacher’s posterior inflates it, so the posterior sees requested-task evidence only.

**Corrections along the way.** Two early claims were retracted on internal review — a teacher-steering read traced to an epoch-frozen sampler (the fixed run confirmed the effect), and a “beats the oracle” comparison traced to a floor handicap in the oracle arm (§6.1 reports the corrected match-within-noise). Full postmortems live in the repository’s audit trail.

## 9 CONCLUSION

For the practical success-conditioned estimator under i.i.d. binary rewards, the expected absolute coefficient mass of a prompt is exactly  $2(\text{pass@}N - \text{pass@}1)$ , and the drop-all-fail implementation targets truncation order  $N-1$ . That identity supplies a compute-aware sampling heuristic and an admission threshold for verified relabels; it does not by itself prove a learning band, optimal prompt allocation, or on-policy hindsight, and we have been explicit about where each stronger reading fails. Empirically, estimator choice and relabel admission changed the trade between mean accuracy and  $\text{pass@}k$  coverage everywhere we looked: an exploratory but zero-exception coverage divergence between MaxRL- and GRPO-labeled maze runs, a three-seed replication of recycling-induced sharpening on Countdown with one gated operating point that buys the coverage back, and a pre-registered negative (Jugs) marking where the whole family of interventions fails together. The confirmatory experiments these results specify — a balanced estimator-by-sampler factorial at six or more seed blocks, a fixed-code gate-strength sweep, and a steering-controlled multi-seed LLM cell — are pre-registered and queued. The portable practice we advocate is independent of our method: measure coverage next to mean accuracy whenever a curriculum or recycler ships, conditioned on the estimator it ships on.

## REFERENCES

- Tajwar, F., Zeng, G., et al. Maximum Likelihood Reinforcement Learning. *ICML*, 2026.
- Li, G., Lin, M., Galanti, T., Tu, Z., Yang, T. DisCO: reinforcing large reasoning models with discriminative constrained optimization. *NeurIPS*, 2025.

- 
- Liu, Z., Cao, Y., Chen, J., Honavar, V. G. Restoring the sweet spot: pass-rate weighted self-distillation for LLM reasoning. *arXiv:2605.27765*, 2026.
- Yue, Y., et al. Does RL really incentivize reasoning capacity in LLMs beyond the base model? *NeurIPS* (oral), 2025.
- Yuan, S., et al. Understanding diversity collapse in RLVR via the lens of overtraining. *arXiv:2606.15455*, 2026.
- Cui, G., et al. The entropy mechanism of RL for reasoning language models. *arXiv:2505.22617*, 2025.
- Cai, W., et al. Curriculum RL can incentivize reasoning capacity beyond the base model. *arXiv:2606.22317*, 2026.
- Wang, Z., Cui, G., Li, Y.-J., Wan, K., Zhao, W. DUMP: automated distribution-level curriculum learning for RL-based LLM post-training. *arXiv:2504.09710*, 2025.
- Shi, T., Wu, Y., Song, L., Zhou, T., Zhao, J. Efficient reinforcement finetuning via adaptive curriculum learning. *TMLR*, 2026. *arXiv:2504.05520*.
- Chen, X., et al. Self-evolving curriculum for LLM reasoning. *arXiv:2505.14970*, 2025.
- Foster, T., Sims, A., Forkel, J., Foerster, J. LILO: learning to reason at the frontier of learnability. *NeurIPS*, 2025.
- Yu, Q., et al. DAPO: an open-source LLM reinforcement learning system at scale. *NeurIPS*, 2025.
- Xiong, W., et al. Reinforce-Ada: an adaptive sampling framework under non-linear RL objectives. *arXiv:2510.04996*, 2025.
- Zheng, H., et al. Act only when it pays: efficient reinforcement learning for LLM reasoning via selective rollouts (GRESO). *NeurIPS*, 2025.
- Andrychowicz, M., et al. Hindsight experience replay. *NeurIPS*, 2017.
- Florensa, C., Held, D., Geng, X., Abbeel, P. Automatic goal generation for reinforcement learning agents. *ICML*, 2018.
- Ren, Z., Dong, K., Zhou, Y., Liu, Q., Peng, J. Exploration via hindsight goal generation. *NeurIPS*, 2019.
- Fang, M., Zhou, T., Du, Y., Han, L., Zhang, Z. Curriculum-guided hindsight experience replay. *NeurIPS*, 2019.
- Pong, V. H., Dalal, M., Lin, S., Nair, A., Bahl, S., Levine, S. Skew-Fit: state-covering self-supervised reinforcement learning. *ICML*, 2020.
- Eysenbach, B., Geng, X., Levine, S., Salakhutdinov, R. Rewriting history with inverse RL: hindsight inference for policy improvement. *NeurIPS*, 2020.
- Ghosh, D., et al. Learning to reach goals via iterated supervised learning. *ICLR*, 2021.
- Ferbach, D., et al. Self-consuming generative models with curated data provably optimize human preferences. *NeurIPS*, 2024.
- Xu, I., Jiang, S., et al. Learning more from less: reinforcement learning from hindsight. *arXiv:2607.09042*, 2026.
- Butt, N., et al. CodeIt: self-improving language models with prioritized hindsight replay. *ICML*, 2024.
- Poesia, G., et al. Learning formal mathematics from intrinsic motivation. *NeurIPS*, 2024.
- Hu, Z., Wang, L., Wang, X., et al. ZipRL: adaptive multi-turn context compression with hindsight response replay. *arXiv:2605.28069*, 2026.
- Strozzi, I. L. Teacher-free self-training amplifies but does not compound: a pass@ $K$  crossover on a free-verifier domain. *arXiv:2606.07856*, 2026.
- Yu, Z., Chen, Z., Liu, M., et al. Unveiling implicit advantage symmetry: why GRPO struggles with exploration and difficulty adaptation. *arXiv:2602.05548*, 2026.
- Rutherford, A., et al. No regrets: investigating and improving regret approximations for curriculum discovery. *NeurIPS*, 2024.




Figure 8: **Gym control under binary task-success reward, trained to plateau (3 seeds).** Solid: mean pass rate across curriculum bins; dashed: the *official* task (MountainCar flag / CartPole survive-400). The target-only arm (gray) is a provable zero-update freeze under sparse verifier-style reward; a uniform-over-bins control (green) also cracks the task,  $3\times$  slower and off-ceiling; the gated FrontierMax stack converges to  $1.000\pm 0.000$ .

## A FULL DETAILS FOR COMPRESSED EXPERIMENT RUNGS

Material moved from §6 during the page-budget pass; every number here is artifact-backed in the repository.

**6.1 control battery (full).** Eight arms, 5 seeds each: uniform 0.650,  $\text{lr}\times 2$  0.798, random-target relabels 0.800, live-gradient replay 0.832, true relabels 0.869, no-floor  $\gamma$ -matched oracle  $0.8885\pm .0014$ , full stack  $0.8895\pm .0023$ , oracle+recycling  $0.8935\pm .0012$ . The relabel direction term (+.037 over replay) has seed spread  $\pm .0025$ , 4–6 $\times$  tighter than every placebo; its worst seed beats every placebo’s best. Random-target relabels are harmful relative to dose-matched replay ( $-.032$ , 5/5 seeds) but neutral vs. plain  $\text{lr}$ -doubling: fake labels waste the slot rather than poison training. Utility-form controls: exact  $u_N$   $0.728\pm .002$  vs. legacy frontier form  $0.733\pm .013$  (the derived form’s payoff is  $7\times$  smaller seed variance); the  $N=2$  slice  $p(1-p)$  is a wash against uniform ( $-.007$ , 2/5 wins) where  $u_N$  carries  $+.078$  (5/5). AUC spread across arms is  $9\times$  the final-performance spread.

**6.5 IsaacLab protocol.** Five arms: frontier teacher, stock greedy walker, uniform, scripted ramp, no-motion control; fixed-grid per-level evaluation; Anymal-C rough terrain; one seed, executed by an independent team against a co-designed, pre-registered protocol. The teacher’s only edge over the greedy walker is a within-noise sliver at the hardest evaluated level.

**6.6 gym control (full).** MountainCar flag: measured random-policy success  $< 3\times 10^{-6}$  over  $10^6$  episodes; survive-400  $\approx 10^{-13}$  by tail fit. Dense-reward controls (prediction registered before the run): CartPole’s native  $+1/\text{step}$  reward lets vanilla REINFORCE on the same tile policy solve survive-400 ( $0.96\pm 0.03$ , 3 seeds); MountainCar’s native  $-1/\text{step}$  reward (constant until first success) stays at 0.000 in all seeds — the categorical zero tracks reward sparsity, not the environment. Attribution at a common 280-step horizon (hard-task AUC): MC  $.750 \rightarrow +\text{recycling} .895 \rightarrow +\text{teacher} .956$ ; CartPole  $.708 \rightarrow .792 \rightarrow .893$ ; paired teacher increment positive 3/3 seeds on both; creation-without-allocation plateaus at 0.958 on CartPole in every seed while adding the teacher restores 1.000; teacher speed effect  $3.0\times$  to hard-task 0.9 ( $2.1\times$  to 0.99 on CartPole). The per-bin ablation replicates the 10-seed transition-matched study (flag  $0.848\pm 0.058$ ) before strengthening it to full convergence.

**6.3 inference-efficiency detail.** Convention: smallest  $k$  reaching absolute coverage 0.25, log-interpolated; artifact `maze.gpu/efficiency.json`; Fig. 9 plots the summed-coverage crossing. Summed coverage over L0–6: GRPO 3.24 vs. MaxRL 3.07 at pass@1; crossing at  $k\approx 4$ ; 5.12 vs. 4.62 at  $k=64$ , with GRPO’s level-2/4 curves saturating at 0.88/0.56 where MaxRL’s reach 1.00/0.62.




Figure 9: **Summed maze coverage vs. inference samples  $k$  (final uniform-arm checkpoints).** GRPO leads at pass@1; the curves cross at  $k \approx 4$ ; MaxRL leads by half a level at  $k=64$  — choose the estimator by how many samples you will draw at inference.

**6.9 secondary reads (full).** Training-side damage: ungated recycling ends with more unsolvable-at-16 frontier training prompts than baseline in every seed (paired  $+ .083 / + .063 / + .034$ ). Entropy at 3 seeds: the gate’s retention replicates ( $+ .072$  paired over baseline, all seeds positive,  $\sim 1.6\times$ ); ungated recycling’s endpoint entropy is seed-unstable (spread  $\pm .062$ ,  $4\text{--}9\times$  the other arms’), consistent with its noisy frontier-coverage endpoints.

**6.7 posterior diagnostics and harness caveats (full).** The pass@ $k$  sweep’s absolute level sits  $\sim 3\times$  below the trainer’s val metric (unreconciled harness discrepancy; only the within-sweep contrast is quoted). Posterior detail: two independent teacher runs’ posteriors agree with each other (cross-run Spearman  $\approx .45$  on co-visited prompts,  $n \approx 900$ ;  $\approx .67$  self-consistency with a shared latent) vs.  $.10\text{--}.22$  against external 7B-model annotations, from  $\sim 1.25$  visits per prompt. Across external-difficulty quintiles spanning solve rates  $.33 \rightarrow .98$ , mean  $\hat{p}$  moves only  $.076 \rightarrow .103$  — the compressed landscape that starves allocation. Response length separates runs by estimator, perfectly: all three GRPO runs compress to  $209.6 \pm 1.8$  tokens while both MaxRL runs sit at  $232.4 \pm 5.7$  (common  $\sim 365$ -token warmstart; the teacher moves length  $\leq 3$  tokens) — sharpening visible in a currency where token entropy is ambiguous. Prompt-uniqueness figures for the replacement confound:  $\sim 34\%$  unique prompts seen by teacher arms vs.  $\sim 43\%$  uniform.

**Jugs suite (pool-conditionality; §8).** Tier grid t0–t4 (2–5 jugs, disjoint min-moves bands, 200 unique tasks/tier); feasibility landscape (SmolLM2-360M two-shot,  $N=16$ ): t0  $.058 / .475$  (pass@1/pass@16), t1  $.006 / .100$ , t2–t4 0/0 with relabel-yield-on-failure  $.62\text{--}.73$  on every tier. All nine pre-registered B-arm cells (no-recycling / ungated / gated  $\times 3$  seeds) converged to the same deterministic fixed point: t0 mean@16 = pass@16 = 0.450 exactly (the 2-move template stratum), t1 pass@16  $0.15 \rightarrow 0$ , actor entropy  $1.36 \rightarrow 0.02\text{--}0.21$ . Ungated recycling ran 169 relabels/step at yield  $.88$  and ended with the lowest entropy ( $0.02\text{--}0.03$ ); the value-keyed gate rejected 99.8% of relabels (the granularity bug described in §8, since fixed and regression-tested). Postmortem and per-cell artifacts in the repository.

## B HYPERPARAMETERS AND COMPUTE

Total compute for every experiment in this paper:  $\sim 230$  A10G-hours on a single GPU plus CPU-days for the exact-gradient suites. Step-time telemetry (artifact generation\_timing.json, 1,188 mixed GSM8K/Countdown steps): step-weighted mean generation fraction 31% overall; the eight generation-dominated sessions ( $>80\%$  session mean, 282 steps) average 83% — a selected-subset description, not an estimate of typical cost. Table 2 gives the per-suite configuration and Table 3 the knob inventory with provenance.



Table 2: Per-suite configuration. Teacher knobs (decay, floor,  $\gamma$ ) are shared across suites with the defaults below; what varies per suite is scale, not schedule.

|                  | chains (§6.1–6.2) | maze (§6.3–6.4)   | GSM8K (§6.7)     | Countdown (§6.8–6.9) |
|------------------|-------------------|-------------------|------------------|----------------------|
| policy           | exact tabular     | 1.26M transformer | SmolLM2-360M     | SmolLM2-360M (SFT)   |
| rollouts $N$     | 16                | 32                | 16               | 16                   |
| groups/step      | 1                 | 8                 | 64               | 64                   |
| optimizer, lr    | SGD, 0.5          | AdamW, $10^{-4}$  | AdamW, $10^{-5}$ | AdamW, $10^{-5}$     |
| budget           | 3,200 groups      | 2,400 s clock     | 50 steps         | 60 steps             |
| seeds (headline) | 5–10              | 3                 | 1–2/cell         | 3/arm                |
| compute/run      | CPU-min           | 0.7 A10G-h        | ~20 A10G-h       | ~6 A10G-h            |

Table 3: Teacher and recycler knobs: value, provenance, and measured sensitivity.

| knob                        | default        | derived or tuned                       | sensitivity                                |
|-----------------------------|----------------|----------------------------------------|--------------------------------------------|
| peak location/ $N$ -scaling | —              | <b>derived</b> (Prop. 1)               | the theory claim                           |
| posterior decay             | 0.7            | tuned (V2b)                            | closes ~19% of oracle gap                  |
| uniform floor               | 0.1            | tuned                                  | 0.15 on maze; robust 0.05–0.2              |
| $\gamma$ (concentration)    | 1              | tuned (V6)                             | 4 on chains; no transfer to maze           |
| gate_max_p                  | 0.5            | $p \rightarrow 1$ zero, midpoint tuned | one trade point validated (§6.9)           |
| relabel dose cap            | 12 groups/step | tuned                                  | sets <i>where</i> coverage is spent (§6.4) |

## REPRODUCIBILITY

Every proposition has an enumeration/Monte-Carlo verification script (`test_mass_formulas.py` in `curriculum_maxrl/` checks the MaxRL, RLOO, and both GRPO-SD mass formulas against exact enumeration and the deployed code). Every figure regenerates from committed inputs listed, with checksums and commands, in the frozen manifest `paper/results/manifest.json`: result values live in versioned JSON tables under `paper/figures/data/` (none are hard-coded in figure scripts), and the maze cohort of Fig. 5 is frozen in an explicit run manifest — adding a run file cannot silently change the figure, and missing inputs are a hard error. The Jugs preregistration, raw cells, verdicts, and postmortem are vendored in `jugs_llm/results/` with preregistration- and execution-commit provenance. Known gaps, stated rather than papered over: the IsaacLab pilot is reported from an independent team’s summary (raw logs live with that team); per-step timing fractions come from verl’s in-run telemetry (`generation_timing.json`); the fig2/fig3 endpoint tables are transcribed from run analyses rather than derived from a seed-level results table, and building that table is queued. The framework is a dependency-light package (`frontier_rl`, numpy-only core) with adapters for LLM pools (verl), gym control, IsaacLab parallel sim, and flow-policy VLAs.